

Machine Learning

Neural Network Overview

- Computers are good at arithmetic but not great at pattern recognition
 - Neural nets attempt to model how neurons transmit information


Chihuahua or Muffin?



(original source unknown)


Neural Network Overview

- Some nomenclature:
 - Neural networks are divided into *layers*
 - There is always an input layer—it doesn't do any processing—just accepts the input
 - There is always an output layer
 - Within a layer, there are neurons or *nodes*
 - For input, there will be one node for each input variable
 - Every node in the first layer connects to every node in the next layer
 - The *weight* associated with the *connection* can vary—these are the matrix elements
 - In this example, the processing is done in layer 2 (output)



Neural Network Overview

- When you train a neural network, you are adjusting the weights connecting the nodes
- Some connections may have zero weight
- This mimics nature—a single neuron can connect to several (or lots) of other neurons
- Linear algebra problem:
 - Inputs: $\mathbf{x} \in \mathbb{R}^n$
 - Outputs: $\mathbf{z} \in \mathbb{R}^m$
 - Neural network is a map, $\mathbb{R}^n \rightarrow \mathbb{R}^m$ that can be expressed as a matrix, \mathbf{A}
 - $\mathbf{z} = \mathbf{Ax}$
 - Given enough input, we could know all the matrix elements in \mathbf{A}



Nonlinear Model

- We'll use a nonlinear function, $g(p)$, that acts on a vector:

$$g(\mathbf{x}) = \begin{pmatrix} g(x_0) \\ g(x_1) \\ \vdots \\ g(x_{n-1}) \end{pmatrix}$$

- then $\mathbf{z} = g(\mathbf{A} \mathbf{x})$
- New procedure: set the entries of \mathbf{A} via training, using a simple, nonlinear, $g(p)$ that fits our training data
- From the graphical representation, the nonlinear function is applied on the output layer


Nonlinear Model

- Again, this mirrors biology
 - Neurons don't act linearly
 - There is a threshold that needs to be reached before a neuron “fires”
- A step function would work, but we want something differentiable
- There are a lot of different choices in the literature

Sigmoid Function

- Common choice: sigmoid function

$$g(p) = \frac{1}{1 + e^{-\alpha p}}$$



- Note, all outputs are scaled to be $z_j \in (0, 1)$
- We'll take $\alpha = 1$

Scaling Output

- Note that since the sigmoid maps all output to (0, 1), we need to make sure that the output in our training set is likewise mapped to (0, 1)
 - If the data doesn't already fall in (0, 1), we can just use a linear transformation:

$$\tilde{x} = 0.9 \frac{x - \min x_i}{\Delta x} + 0.05$$

- here, Δx is the largest possible range of x_i in the inputs

Actually, (0, 1] works fine—we just need to avoid a 0, since that cancels out weights in the matrices

Implementation

- Basic operation
 - Train the model with known input/output to get all A_{ij}
 - Use $\mathbf{z} = g(\mathbf{A} \mathbf{x})$ to get output for a new input \mathbf{x}
- Training:
 - We have T pairs $(\mathbf{x}^k, \mathbf{y}^k)$ for $k = 1, \dots, T$
 - Important: remember that our \mathbf{y} 's have to be scaled to be in $(0, 1)$, so they are in the same range that our function $g(p)$ maps to
 - We require that $g(\mathbf{A} \mathbf{x}^k) = \mathbf{y}^k$ for all k
 - Recall, that $g(p)$ is a scalar function that works element-by-element:

$$z_i = g([\mathbf{A}\mathbf{x}]_i) = g\left(\left[\sum_j A_{ij}x_j\right]\right)$$

Implementation

- Training (cont.)
 - We find the elements of \mathbf{A}
 - This can be expressed as a minimization problem, where we alter the matrix elements to achieve this agreement
 - There may not be a unique set of A_{ij} , so we will loop randomly over all training data multiple times to optimize \mathbf{A}

$$f(A_{ij}) = \|g(\mathbf{A}\mathbf{x}^k) - \mathbf{y}^k\|^2$$

- This looks like a least-squares minimization
- The function we minimize is called the *cost function*
 - There are other choices than the square of the error

Implementation

- Minimization
 - A common technique for minimization is *gradient descent* (sometimes called steepest descent)
 - This looks at the local derivative of the function f with respect to the parameters A_{ij} and moves a small distance *downhill*, and iterates...
 - We'll also compare to an external library for minimization
- Caveats
 - When you minimize with one set of training data, there is no guarantee that you are still minimized with respect to the previous sets
 - In practice, you feed the training data multiple times, in random order, to the minimizer—each pass is called an *epoch*

Aside: Minimization

- Steepest descent minimization
 - Start at a point \mathbf{x}_0 and evaluate the gradient
 - Move *downhill* by following the gradient by some amount η
 - Correct our initial guess and iterate

$$\mathbf{x} \leftarrow \mathbf{x} - \eta \frac{\partial u}{\partial \mathbf{x}}$$


- Need to choose the amount to move each iteration
 - Sometimes we instead define a unit vector in the direction of the local gradient, and then η represents the distance to travel in that direction
- You can think about this as what happens if you put a marble on a surface—it rolls to a minimum
 - May not be the global minimum—we can get stuck in a local minimum

Aside: Minimization

- Example: Rosenbrock (banana) function

$$f(x, y) = (a - x)^2 + b(y - x^2)^2$$


- This is a hard problem for optimization
- Global minimum is at (a, a^2)



Note: this is the
log of the
function plotted

Aside: Minimization

- Minimization with gradient descent is very sensitive to choice of η
 - Too large and you may shoot off far from the minimum
 - Too small and you do a lot of extra work



Neural Net Minimization

- For our function,

$$f(A_{ij}) = \|g(\mathbf{A}\mathbf{x}^k) - \mathbf{y}^k\|^2$$

- Note, this definition is for a single training pair, $(\mathbf{x}^k, \mathbf{y}^k)$

$$(\mathbf{x}^k, \mathbf{y}^k) = (\{x_1^k, x_2^k, \dots, x_n^k\}, \{y_1^k, y_2^k, \dots, y_m^k\})$$

- Our update would be

$$A_{pq} = A_{pq} - \eta \frac{\partial f}{\partial A_{pq}}$$

- where

$$f(A_{ij}) = \sum_{i=1}^m \left[g \left(\sum_{j=1}^n A_{ij} x_j \right) - y_i \right]^2$$

Neural Net Minimization

- Working out the derivative:

$$\frac{\partial f}{\partial A_{pq}} = 2(z_p - y_p)\alpha z_p(1 - z_p)x_q$$

- We could then use steepest descent, looping over the matrix elements and doing the minimization on them one by one, iterating until we converge
 - Instead, we just do one push “downhill” following the gradient for a single training set and then move to the next.
 - η is often called the *learning rate*
- Gradient descent is often used for neural nets because it only requires the first derivative
 - Newton’s method would require the second derivatives (Hessian matrix)

Neural Net Minimization

- Recall,
 - \mathbf{A} is $m \times n$ matrix
 - \mathbf{x} is $n \times 1$ vector
 - \mathbf{y} (and hence \mathbf{z}) is $m \times 1$ vector
- We can write our derivative as:

$$\frac{\partial f}{\partial \mathbf{A}} = \underbrace{2(\mathbf{z} - \mathbf{y}) \circ \alpha \mathbf{z} \circ (1 - \mathbf{z})}_{m \times 1} \cdot \underbrace{\mathbf{x}^T}_{1 \times n}$$


- Then the correction to our matrix is:

$$\begin{aligned}\Delta \mathbf{A} &= -2\eta (\mathbf{z} - \mathbf{y}) \circ \alpha \mathbf{z} \circ (1 - \mathbf{z}) \cdot \mathbf{x}^T \\ \mathbf{A} &\leftarrow \mathbf{A} + \Delta \mathbf{A}\end{aligned}$$

Here, $\mathbf{a} \circ \mathbf{b}$ is an
element-wise
product

Hidden Layers

- We can add more more parameters by another layer of nodes/neurons



Hidden Layers


- *Hidden layers* sit between the input and output
- For hidden layer of dimension k :
 - Inputs: $\mathbf{x} \in \mathbb{R}^n$
 - Outputs: $\mathbf{z} \in \mathbb{R}^m$
 - \mathbf{A} is an $m \times k$ matrix
 - \mathbf{B} is an $k \times n$ matrix
 - The product \mathbf{AB} is $m \times n$, as we had before
- *Universal approximation theorem*: single layer network can represent any continuous function
- *We transform the input in two steps:*

$$\tilde{\mathbf{z}} = g(\mathbf{B}\mathbf{x})$$

$$\mathbf{z} = g(\mathbf{A}\tilde{\mathbf{z}})$$

Hidden Layers

- Graphically this appears as:



Hidden Layers

- Now we minimize:

$$f(A_{ls}, B_{ij}) = \sum_{l=1}^m (z_l - y_l)^2$$
$$\tilde{z}_i = g \left(\sum_{j=1}^n B_{ij} x_j \right)$$
$$z_l = g \left(\sum_{s=1}^k A_{ls} \tilde{z}_s \right)$$

Minimization

- We need to do the minimization now for both sets of weights (matrices)
- In practice, we do them one at a time, with each seeing the result from its layer
 - This process is also called *backpropagation* in neural networks—we are using the errors at the end to change the weights that came earlier in the network

Gradient Descent

- We can do our gradient descent on **A** and **B** separately now
 - This is the strength of backpropagation and gradient descent vs. some “canned” minimization routine—we are not optimizing the entire system all together
- Differentiating our error and lots of chain rule gives:

$$\Delta A = -2\eta \mathbf{e} \circ \mathbf{z} \circ (1 - \mathbf{z}) \cdot \tilde{\mathbf{z}}^\top$$

$$\Delta B = -2\eta \tilde{\mathbf{e}} \circ \tilde{\mathbf{z}} \circ (1 - \tilde{\mathbf{z}}) \cdot \mathbf{x}^\top$$

Note: this is a single dot product, the combination of vectors on the left are multiplied element-by-element (the Hadamard product)

- With

$$\tilde{\mathbf{e}} = \mathbf{A}^\top \mathbf{e} \circ \mathbf{z} \circ (1 - \mathbf{z}) \approx \mathbf{A}^\top \mathbf{e}$$


This approximation seems to be commonly made and supposedly doesn't affect convergence much

Image Classification

- We'll try to recognize a digit (0 – 9) from an image of a handwritten digit.
 - MNIST dataset (<http://yann.lecun.com/exdb/mnist/>)
 - Popular dataset for testing out machine learning techniques
 - Training set is 60,000 images
 - Approximately 250 different writers
 - Test set is 10,000 images
 - Correct answer is known for both sets so we can test our performance
- Image details:
 - 28×28 pixels, grayscale (0 – 255 intensity)
- The best learning algorithms can get accuracy > 99%

Image Classification

- Neural network characteristics:
 - Input layer will be 784 nodes
 - One for each pixel in the input image
 - Output layer will be 10 nodes
 - An array with an entry for each possible digit
 - “3” would be represented as: $[0, 0, 0, 1, 0, 0, 0, 0, 0, 0]$
 - We’ll start with a hidden layer size of 100
- We’ll train on the training set, using up to 60000 images
 - Rescale the input to be in $[0.01, 1]$
- We’ll test on the test set of 10000 images





First digit MNIST in the training set



Image Classification

- Default configuration:
 - The full training set (60000 images)
 - Hidden layer of 100 nodes
 - 5 epochs of training
 - Learning rate of 0.1
- We achieve 95 – 96% accuracy

Some Image Classification Failures



Some Image Classification Failures



Some Image Classification Failures








Image Classification Weights

- Weights (matrix elements of **A** and **B**) seem symmetric about 0
 - Interestingly, with more training, the width of the distribution seems to grow




Effect of Number of Epochs

- When we use the full training set (60000 images) the number of epochs (passes through the training data) doesn't seem to matter much




Effect of Training Set Size

- No surprise: the larger the training set, the better we do



Effect of Hidden Layer Size

- Also not unexpected: the larger the hidden layer the better we do



Effect of Learning Rate

- A smaller learning rate seems to do better

